

04-630

Data Structures and Algorithms for Engineers

Lecture 15: Priority Queues

Agenda

- Priority queues
- Binary heap
- Applications
 - Heapsort

Priority Queues

- Many applications require algorithms to process items in a specific order (e.g., relative importance).
 - One option: Use a list, **sort it**, and process in the resultant order(ascending or descending)
- Priority queues are more flexible
 - They allow new elements to be **added at arbitrary intervals**
 - **More efficient** to add the new element in a priority queue than to **add** to a list and **re-sort**

Priority Queues

Main priority queue operations

- *Insert* (Q, x)

Given an item x with a key k , insert it into the priority queue Q

- Find_Minimum(Q) or Find_Maximum(Q)

Return a pointer to the item whose key value is smaller / larger than any other key in the priority queue

- Delete_Minimum(Q) or DeleteMaximum(Q)

Remove the item from the priority queue Q whose key is minimum or maximum

Priority Queues

Possible implementations

- Unsorted array
- *Sorted array*: inserting new elements is slow.
- Balanced binary search tree
- *Binary heap*: suitable when you know upper bound of elements in priority queue (since size of array needs to be specified upfront)....but can mitigate this too using dynamic arrays.

Priority Queues

Possible implementations

	Unsorted array	Sorted array	Balanced tree
Insert(Q, x)	$O(1)$	$O(n)$	$O(\log n)$
Find-Minimum(Q)	$O(1)$	$O(1)$	$O(1)$
Delete-Minimum(Q)	$O(n)$	$O(1)$	$O(\log n)$

- How do we get $O(1)$ for Find-Minimum(Q) in all three cases?
- Use a variable to store a pointer/index to the minimum entry in each of these structures
 - Update the pointer on each insertion (if necessary)
 - Update it on each deletion, requiring the new minimum to be found (but the cost of this can be folded into the cost of the deletion because it is the same cost as the deletion)

Binary Heap

- Favoured implementation of priority queue
- Heaps maintain **partial order** on the set of elements
 - Weaker than sorted order (& so **it is efficient**)
 - Stronger than random order (& so min/max element can be **quickly identified**)
 - “Heap” refers to being “**top of the heap**”, i.e., what’s on top dominates what is underneath
 - **Dominate**: is **greater than** (or **less than**) everything under it

Binary Heap

Heap-labelled tree is a binary tree

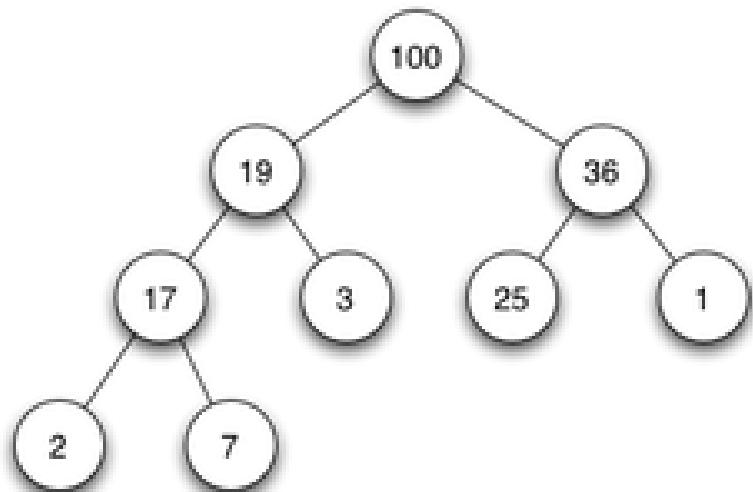
- Keys in each node **dominate** the keys of each of its children
- **Min-heap** ... less than its children
- **Max-heap** ... greater than its children

Binary Heap

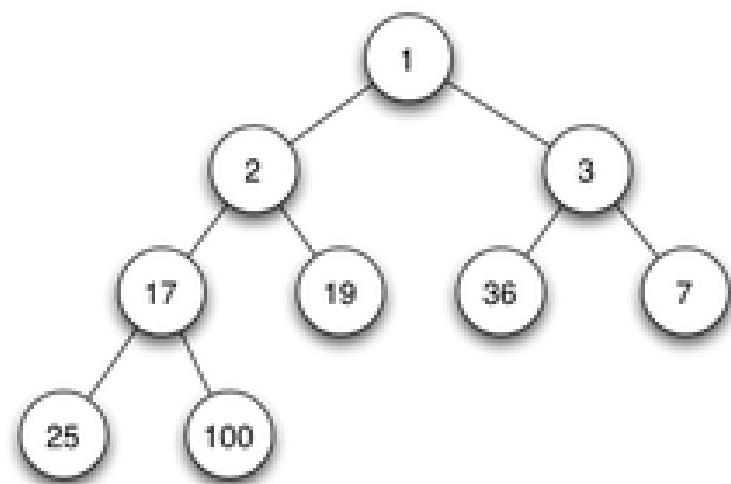
A ***binary heap*** is a ***binary tree*** that satisfies **two special shape and heap properties:**

- all levels of the tree, except possibly the last one (deepest) are **fully filled**, and, if the last level of the tree is not complete, the nodes of that level are filled from **left to right**
- each node is “**greater than or equal to**” each of its **children** (in the case of a **max-heap**) according to some **comparison predicate** which is fixed for the entire data structure

Binary Heap



Max-heap



Min-heap

Binary Heap

The order of siblings is not specified

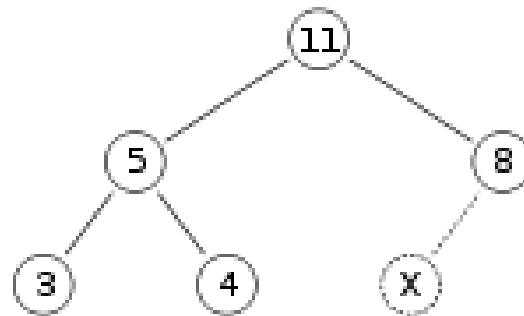
- Two children can be freely interchanged
- As long as it **doesn't violate** the **shape and heap** properties

Binary Heap

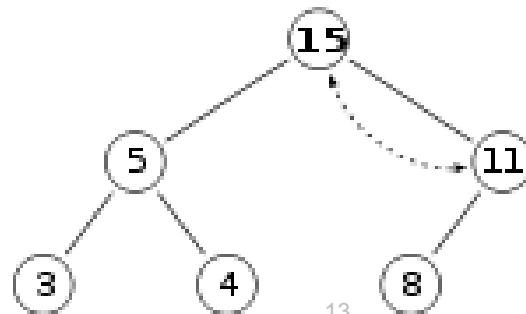
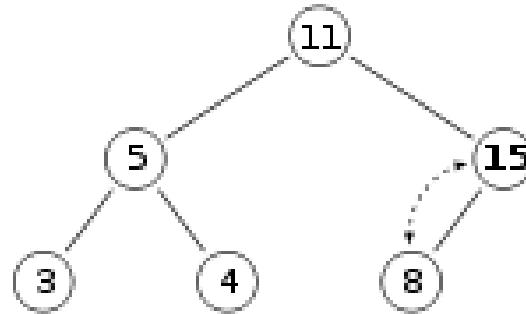
Adding to a heap

- Algorithm: **upheap / heapify-up / sift-up**
 - Add element to bottom level
 - Compare the added element with its parent;
if they are in correct order, stop
 - If not, swap the element with its parent and return to previous step
- $O(\log n)$

Binary Heap



Add 15 to max-heap



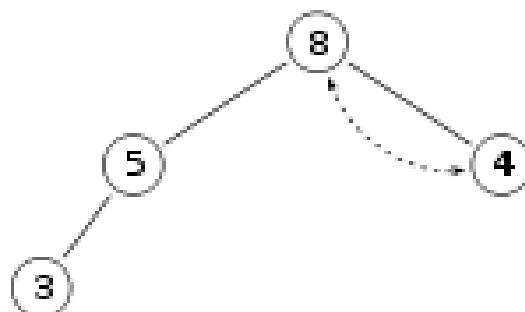
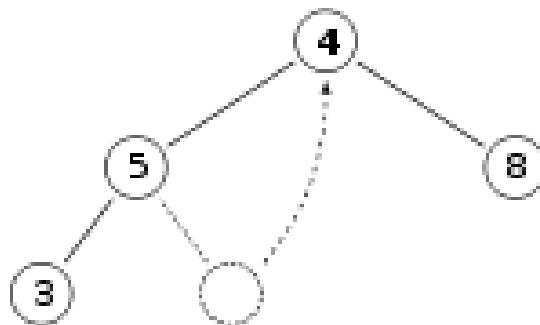
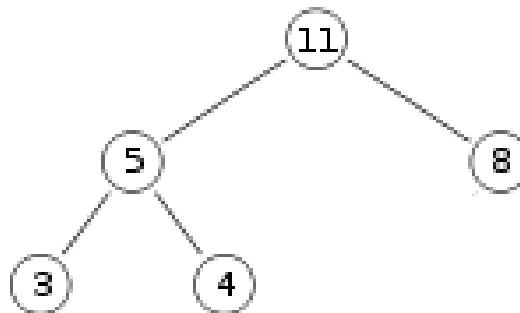
Binary Heap

Deleting the root from a heap

- Effectively extracting the **maximum element** in a **max-heap** (or extracting the **minimum element** in **min-heap**)
- Algorithm: **downheap / heapify-down / sift-down**
 - Replace root with last element on the bottom level
 - Compare the swapped element with
 - The larger child (max-heap)
 - The smaller child (min-heap)
 - if they are in correct order, stop
 - If not, swap the element with the child and return to previous step
- $O(\log n)$

Binary Heap

Delete Root from max-heap



Binary Heap

Implementation as a binary tree data structure

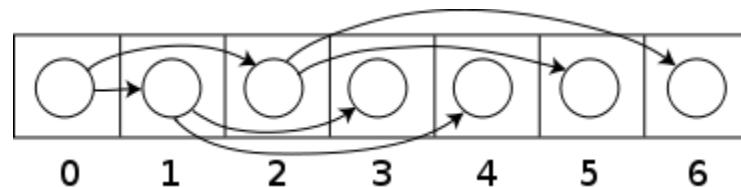
Problem: finding adjacent element on last level

- Find it algorithmically (takes time)
- Use additional links between siblings:
threading the tree (takes space)

Binary Heap

Implementation as an array

- Represent a binary tree WITHOUT any pointers by using an array of keys and a **mapping function**
- Use formula to find parents and children of a node
 - Node at index i has children at indices $2i + 1$ and $2i + 2$
 - Node at index i has parent at index $(i - 1)/2$



Binary Heap

```
/* Implementation of a min-heap priority queue */  
/*  
 * To simplify matters, assume the array starts at index 1 */  
 * Thus, left child of key at index k is located at index 2k */  
 * and right child is located at index 2k+1 */  
 * and parent of key at index k is located at index floor(k/2) */  
/*  
 * Exercise: rewrite the code so that array starts at index 0 */
```

```
typedef struct {  
    item_type q[PQ_SIZE+1];           /* body of queue */  
    int n;                            /* number of queue elements */  
} priority_queue;
```

Binary Heap

```
pq_parent(int n){  
    if (n == 1) return(-1);  
    else return((int) n/2);      /* implicitly take floor(n/2) */  
}  
  
pq_young_child(int n)  
{  
    return(2 * n);  
}
```

Binary Heap

```
/* Construct the heap incrementally */  
/*  
 * insert a new element at the left-most open location in the array */  
 * i.e. at index (n+1) of a previously n-element heap.  
 */  
/*  
 * This guarantees the desired balance but not necessarily the */  
 * dominance ordering of the keys: new key might be less than its */  
 * parent in a min-heap or greater than its parent in a max-heap */  
/*  
 * To satisfy this condition, we recursively bubble up the new key */  
 * to its correct position (i.e. upheap / heapify-up / sift-up) */  
/*  
 * Since the heap binary tree is almost full, its height is log n */  
 * and consequently each insertion takes at most O(log n) time */  
 * Thus an initial heap of n elements can be constructed in */  
 * O(n log n) time through n insertions */
```

Binary Heap

```
 pq_insert(priority_queue *q, item_type x) {
    if (q->n >= PQ_SIZE)
        printf("Warning: priority queue overflow insert x=%d\n", x);
    else {
        q->n = (q->n) + 1; //first item will be inserted at point
                               // (0+1)=1
        q->q[ q->n ] = x; //shape condition
        bubble_up(q, q->n); //dominance condition
    }
}

bubble_up(priority_queue *q, int p) {
    if (pq_parent(p) == -1) return; /* at root of heap, no parent */
    if (q->q[pq_parent(p)] > q->q[p]) { //remember : min-heap--parent
                                              //greater than child
        pq_swap(q,p,pq_parent(p)); //queue and the two indices to
                                     //be swapped
        bubble_up(q, pq_parent(p));
    }
}
```

Binary Heap

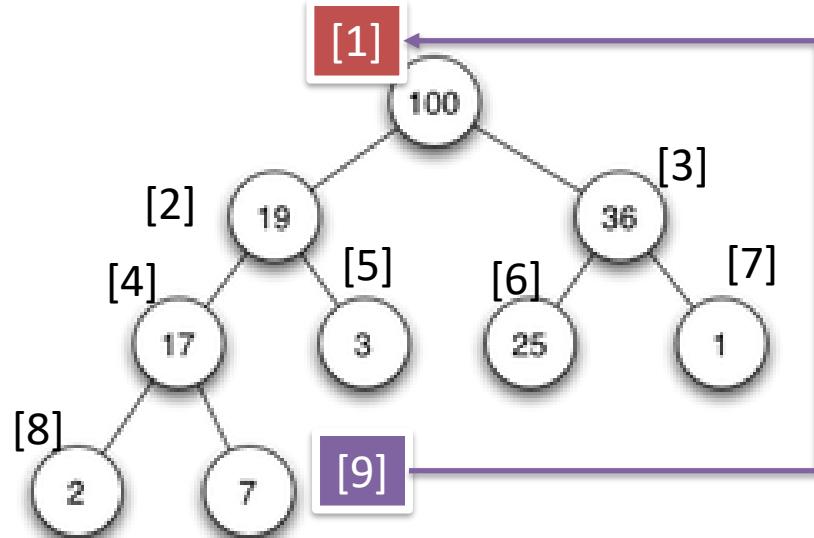
```
 pq_init(priority_queue *q) {  
     q->n = 0;  
 }  
  
 make_heap(priority_queue *q, item_type s[], int n)  
     int i; /* counter */  
  
     pq_init(q);  
  
     for (i=0; i<n; i++)  
         pq_insert(q, s[i]);  
 }  
  
 pq_swap(priority_queue *q, int i, int j)  
 {  
     int temp;//placeholder  
     temp = q->q[i];  
     q->q[i] = q->q[j];  
     q->q[j] = temp;  
 }  
 
```

Binary Heap

```
/* Extract the minimum */  
/*  
/* Identify and delete the dominant element: */  
/* i.e. the element at the root of the tree, i.e the first element */  
/*  
/* Deleting this root element leaves a hole in the array */  
/* which we fill by moving the element from the right-most leaf */  
/* (at the nth position of the array) into the first position */  
/*  
/* Again, this satisfies the shape condition but not necessarily */  
/* the dominance condition. */  
/* So, we recursively bubble down the new root */  
/* to its correct position (i.e. heapify) */  
/*  
/* Root deletion is achieved in O(log n) time */
```

Deleting and moving last leaf item: max- & min-heaps

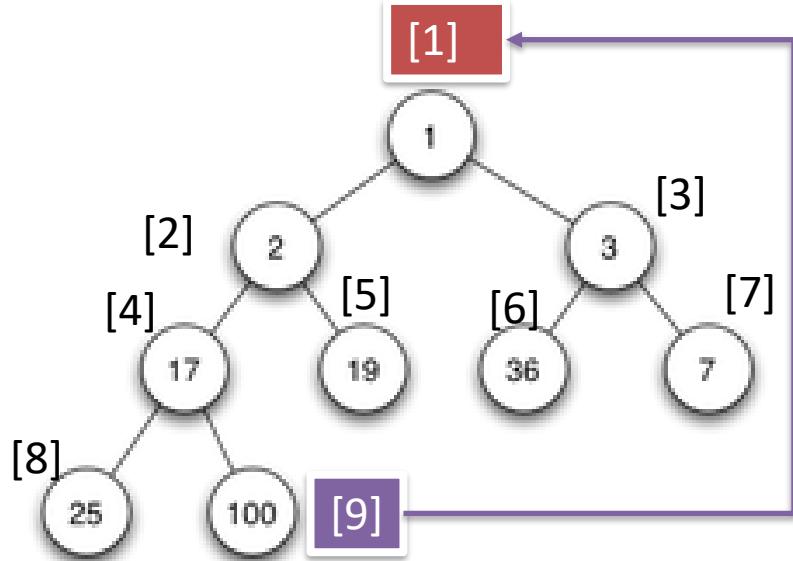
Delete 100, move 7, bubble down



[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]
100	19	36	17	3	25	1	2	7

Max-heap

Delete 1, move 100, bubble down



[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]
1	2	3	17	19	36	7	25	100

Min-heap

Recall: for index k , left_child_index=2k, right_child_index=2k+1

Binary Heap

```
item_type extract_min(priority_queue *q) {  
  
    int min = -1;                                /* minimum value */  
  
    if (q->n <= 0)  
        printf("Warning: empty priority queue.\n");  
    else {  
        min = q->q[1];//root  
  
        q->q[1] = q->q[ q->n ];//filling hole in array  
        q->n = q->n - 1;//reduce number of elements in array-  
                           //shape condition  
        bubble_down(q,1);//dominance condition/property  
    }  
  
    return(min);  
}
```

Binary Heap

```
bubble_down(priority_queue *q, int p) {  
  
    int c;          /* child index */  
    int i;          /* counter */  
    int min_index; /* index of lightest child */  
  
    c = pq_young_child(p); //recall pq_young_child returns 2*p  
  
    min_index = p;  
  
    /* find the index of the smallest of parent and children */  
    for (i=0; i<=1; i++)  
        if ((c+i) <= q->n) { // check that the child is there  
            if (q->q[min_index] > q->q[c+i])  
                min_index = c+i;  
        }  
  
    if (min_index != p) { // if the smallest of the three nodes  
        pq_swap(q,p,min_index); // is NOT the parent, we need to swap  
        bubble_down(q, min_index); // and repeat with that child  
    }  
}
```

Applications of Priority Queues

- Implementation of:
 - Sorting data, e.g., heapsort
 - Dijkstra's shortest path algorithms.
 - Huffman coding used in data compression.
 - minimum spanning tree algorithms e.g., Prim's algorithm.
 - solutions to fetch the next best or next worst element.
 - best first search algorithms e.g., A* search (fetches next best).

Example application: Heapsort

- The heap priority queue can be used to create a very efficient sorting algorithm: heapsort
- Construct the priority queue: $O(n \log n)$
- Repeatedly extract the minimum: $O(n \log n)$
- Overall complexity is $O(n \log n)$ worst-case
- This is the best that can be expected from any sorting algorithm
- Also, it is an in-place sort, meaning it uses no extra memory in addition the array containing the elements is to be sorted

Heapsort

```
heapsort(item_type s[], int n) {  
  
    int i;          /* counters */  
    priority_queue q; /* heap for heapsort */  
  
    make_heap(&q,s,n); //populate priority queue  
  
    for (i=0; i<n; i++)  
        s[i] = extract_min(&q); //repeatedly extract the root  
}
```

Heap vs BST and AVL

Heap Tree				BST	AVL
	Unsorted array	Sorted array	Balanced tree		
Insert(Q, x)	$O(1)$	$O(n)$	$O(\log n)$	$O(n)$	$O(\log n)$
Find-Minimum(Q)	$O(1)$	$O(1)$	$O(1)$	$O(n)$ -search	$O(\log n)$ -search
Delete-Minimum(Q)	$O(n)$	$O(1)$	$O(\log n)$	$O(n)$ - delete	$O(\log n)$ - delete

Acknowledgement

Adopted and Adapted from Material by:

David Vernon: www.vernon.eu

Augmented by material from:

The Algorithm Design Manual 2nd Edition: by Steven Skiena

Introduction to Algorithms, 3rd Edition, Thomas H. Cormen et al. (2009)